

# ADL21-HW2

B08303028 Hsiang-An Chuang

November 2022

## 1 Model

1. The model we use in this chinese summarization task is `google/mt5-small`. `mt5-small` is a encoder-decoder Transformer which is a multilingual seq2seq pre-trained model. For every input tokens, they go through encoder, which is stacks of blocks. In each block, there is a self-attention layer followed by a feed-forward neural network, and there is a layer normalization at the end of each block. As for decoder, the only difference is that there is a standard attention mechanism after each self-attention layer and it feed the output of the final block with a linear layer to get the outputs. Thus, the generation strategies can operated on these outputs
2. As for preprocessing, mostly I apply `preprocess_function` in `run_summarization.py` provided by huggingface.

```
1 def preprocess_function(examples):
2     # remove pairs where at least one record is None
3     prefix = ""
4     padding = 'max_length'
5     max_target_length = args.max_target_length
6
7     inputs, targets = examples['maintext'], examples['
8     title']
9
10    inputs = [prefix + inp for inp in inputs]
11    model_inputs = tokenizer(inputs, max_length=args.
12    max_source_length, padding=padding, truncation=True)
13
14    # Tokenize targets with the `text_target` keyword
15    argument
16    labels = tokenizer(text_target=targets, max_length=
17    max_target_length, padding=padding, truncation=True)
18
19    # If we are padding here, replace all tokenizer.
20    pad_token_id in the labels by -100 when we want to ignore
21    # padding in the loss.
22    if padding == "max_length":
23        labels["input_ids"] = [
24            [(1 if l != tokenizer.pad_token_id else -100)
25             for l in label] for label in labels["input_ids"]]
```

```

20         ]
21
22         model_inputs["labels"] = labels["input_ids"]
23         return model_inputs
24

```

The algorithm simply put maintext, summary(title in our task) together and add a prefix in each text. In term of maintext, it pad all maintext to `max_input_length`, which is 512 by default. As for summary, they are padded to `max_target_length`, which i had tried 64/256 in differencnt experiments.

## 2 Training


1. The hyperparameters I chose mostly consider the problem of 'cuda out of memory'.

```

1  {
2      "num_train_epochs": 30
3      "per_device_train_batch_size": 2
4      "per_device_eval_batch_size": 2
5      "max_source_length": 512
6      "max_target_length": 64
7      "gradient_accumulation_steps": 16
8      "fp16": True
9  }

```

2. Plot the learning curves (ROUGE versus training steps)



(rouge-1, rouge-2, rouge-l): (24.98, 9.23, 22.28)

### 3 Generation Strategies

#### 1. Strategies

- **Greedy:** Choose the most probable word for each output, and not consider the global maximum.
- **Beam Search:** Keep track of the k-most probable sequences and finding a better one.
- **Top-k Sampling:** Kind of sampling, sample the word via distribution but restricted to the top-k probable words.
- **Top-p Sampling:** Kind of sampling, sample from a subset of vocabulary with the most probability mass.
- **Temperature:** Apply hyperparameter  $\tau$  in softmax function. Set a higher value to smooth the distribution.

#### 2. Hyperparameters

| Strategy                                  | rouge-1      | rouge-2      | rouge-l      |
|---|--------------|--------------|--------------|
| <b>Greedy</b>                             | 0.244        | 0.091        | 0.232        |
| <b>Beam Search(beam=3)</b>                | <b>0.267</b> | 0.104        | 0.238        |
| <b>Beam Search(beam=5)</b>                | 0.266        | <b>0.106</b> | <b>0.239</b> |
| <b>Top-k Sampling(k=10)</b>               | 0.231        | 0.068        | 0.193        |
| <b>Top-k Sampling(k=20)</b>               | 0.219        | 0.064        | 0.184        |
| <b>Top-p Sampling(p=0.8)</b>              | 0.187        | 0.053        | 0.162        |
| <b>Top-p Sampling(p=0.9)</b>              | 0.175        | 0.051        | 0.151        |
| <b>Temperature(<math>\tau=0.5</math>)</b> | 0.249        | 0.081        | 0.210        |
| <b>Temperature(<math>\tau=0.8</math>)</b> | 0.249        | 0.081        | 0.210        |

My final generation strategy is **Beam Search(beam=5)**.